

Acceleration Structures for Ray Casting

A 3D rendered scene of an ancient Egyptian temple interior. The scene features several large, fluted columns with papyrus capitals. The walls are covered in hieroglyphs and have decorative moldings. The floor is made of large stone tiles. The lighting is warm, suggesting an interior space with light coming from above.

MIT EECS 6.837 Computer Graphics
Wojciech Matusik, MIT EECS

Recap: Ray Tracing

trace ray

Intersect all objects

color = ambient term

For every light

 cast shadow ray

 color += local shading term

If mirror

 color += color_{refl} *


 trace reflected ray

If transparent

 color += color_{trans} *

 trace transmitted ray


- *Does it ever end?*




Stopping criteria:

- Recursion depth
 - Stop after a number of bounces
- Ray contribution
 - Stop if reflected / transmitted contribution becomes too small


Recursion For Reflection: None



Recursion For Reflection: 1




Recursion For Reflection: 2



Ray tree


- Visualizing the ray tree for single image pixel



Ray tree

This gets pretty complicated
pretty fast!

- Visualizing the ray tree for single image pixel



Questions?

Ray Tracing Algorithm Analysis

- Lots of primitives
- Recursive
- Distributed Ray Tracing
 - Means using many rays for non-ideal/non-pointlike phenomena
 - Soft shadows
 - Anti-aliasing
 - Glossy reflection
 - Motion blur
 - Depth of field

cost \approx height * width *
num primitives *
intersection cost *
size of recursive ray tree *
num shadow rays *
num supersamples *
num glossy rays *
num temporal samples *
num aperture samples *
...

Can we reduce this?

Today

- Motivation
 - You need LOTS of rays to generate nice pictures
 - Intersecting every ray with every primitive becomes the bottleneck
- Bounding volumes
- Bounding Volume Hierarchies, Kd-trees

For every pixel

Construct a ray from the eye

For every object in the scene


Find intersection with the ray

Keep if closest

Shade


Accelerating Ray Casting

- Goal: Reduce the number of ray/primitive intersections




Conservative Bounding Volume

- First check for an intersection with a conservative bounding volume
- Early reject: If ray doesn't hit volume, it doesn't hit the triangles!






Conservative Bounding Volume

- What does “conservative” mean?
 - Volume must be big enough to contain all geometry within




Conservative Bounding Regions

- Desiderata
 - Tight → avoid false positives
 - Fast to intersect




Ray-Box Intersection

- Axis-aligned box
- Box: $(X_1, Y_1, Z_1) \rightarrow (X_2, Y_2, Z_2)$
- Ray: $P(t) = R_o + tR_d$




Naïve Ray-Box Intersection

- 6 plane equations: Compute all intersections
- Return closest intersection *inside the box*
 - Verify intersections are on the correct side of each plane: $Ax+By+Cz+D < 0$




Reducing Total Computation

- Pairs of planes have the same normal
- Normals have only one non-zero component
- Do computations one dimension at a time



Test if Parallel


- If $R_{dx} = 0$ (ray is parallel) AND $R_{ox} < X_1$ or $R_{ox} > X_2 \rightarrow$ **no intersection**



**(The same
for Y and Z,
of course)**


Find Intersections Per Dimension

- Basic idea
 - Determine an interval along the ray for each dimension
 - The intersect these 1D intervals (remember CSG!)
 - Done!




Find Intersections Per Dimension

- Basic idea
 - Determine an interval along the ray for each dimension
 - The intersect these 1D intervals (remember CSG!)
 - Done!



Find Intersections Per Dimension

- Basic idea
 - Determine an interval along the ray for each dimension
 - The intersect these 1D intervals (remember CSG!)
 - Done!




Interval
between X_1
and X_2

Interval
between Y_1
and Y_2

Find Intersections Per Dimension

- Basic idea
 - Determine an interval along the ray for each dimension
 - The intersect these 1D intervals (remember CSG!)
 - Done!




Interval
between X_1
and X_2

Interval
between Y_1
and Y_2


Intersection

Intersecting 1D Intervals




Intersecting 1D Intervals

Start=
max of mins



Intersecting 1D Intervals


Start=
max of mins



End=
min of maxs


Intersecting 1D Intervals

If $\text{Start} > \text{End}$, the intersection is empty!




Find Intersections Per Dimension

- Calculate intersection distance t_1 and t_2




Find Intersections Per Dimension

- Calculate intersection distance t_1 and t_2
 - $t_1 = (X_1 - R_{ox}) / R_{dx}$
 - $t_2 = (X_2 - R_{ox}) / R_{dx}$
 - $[t_1, t_2]$ is the X interval




Then Intersect Intervals

- Init t_{start} & t_{end} with X interval
- Update t_{start} & t_{end} for each subsequent dimension




Then Intersect Intervals

- Compute t_1 and t_2 for Y...




Then Intersect Intervals

- Update t_{start} & t_{end} for each subsequent dimension
 - If $t_1 > t_{\text{start}}$, $t_{\text{start}} = t_1$
 - If $t_2 < t_{\text{end}}$, $t_{\text{end}} = t_2$




Then Intersect Intervals

- Update t_{start} & t_{end} for each subsequent dimension
 - If $t_1 > t_{\text{start}}$, $t_{\text{start}} = t_1$
 - If $t_2 < t_{\text{end}}$, $t_{\text{end}} = t_2$



Then Intersect Intervals


- Update t_{start} & t_{end} for each subsequent dimension
 - If $t_1 > t_{\text{start}}$, $t_{\text{start}} = t_1$
 - If $t_2 < t_{\text{end}}$, $t_{\text{end}} = t_2$



:-)


Is there an Intersection?

- If $t_{\text{start}} > t_{\text{end}}$ → **box is missed**




Is the Box Behind the Eyepoint?

- If $t_{\text{end}} < t_{\text{min}}$ → **box is behind**



Return the Correct Intersection

- If $t_{\text{start}} > t_{\text{min}}$ → **closest intersection at t_{start}**
- Else → **closest intersection at t_{end}**
 - Eye is inside box




Ray-Box Intersection Summary

- For each dimension,
 - If $R_{dx} = 0$ (ray is parallel) AND $R_{ox} < X_1$ or $R_{ox} > X_2 \rightarrow$ **no intersection**
- For each dimension, calculate intersection distances t_1 and t_2
 - $t_1 = (X_1 - R_{ox}) / R_{dx}$ $t_2 = (X_2 - R_{ox}) / R_{dx}$
 - If $t_1 > t_2$, swap
 - Maintain an interval $[t_{start}, t_{end}]$, intersect with current dimension
 - If $t_1 > t_{start}$, $t_{start} = t_1$ If $t_2 < t_{end}$, $t_{end} = t_2$
- If $t_{start} > t_{end} \rightarrow$ **box is missed**
- If $t_{end} < t_{min} \rightarrow$ **box is behind**
- If $t_{start} > t_{min} \rightarrow$ **closest intersection at t_{start}**
- Else \rightarrow **closest intersection at t_{end}**


Efficiency Issues

- $1/R_{dx}$, $1/R_{dy}$ and $1/R_{dz}$ can be pre-computed and shared for many boxes

Bounding Box of a Triangle




Bounding Box of a Sphere



Bounding Box of a Plane


$$(x_{max}, y_{max}, z_{max}) \\ = (+\infty, +\infty, +\infty)^*$$



$$(x_{min}, y_{min}, z_{min}) \\ = (-\infty, -\infty, -\infty)^*$$

** unless n is exactly perpendicular to an axis*

Bounding Box of a Group




$$\begin{aligned} &= (\max(x_{\max_a}, x_{\max_b}), \\ &\quad \max(y_{\max_a}, y_{\max_b}), \\ &\quad \max(z_{\max_a}, z_{\max_b})) \end{aligned}$$

$$\begin{aligned} (x_{\min}, y_{\min}, z_{\min}) &= (\min(x_{\min_a}, x_{\min_b}), \\ &\quad \min(y_{\min_a}, y_{\min_b}), \\ &\quad \min(z_{\min_a}, z_{\min_b})) \end{aligned}$$


Bounding Box of a Transform

Bounding box of transformed object IS NOT the transformation of the bounding box!




Bounding Box of a Transform

Bounding box of transformed object IS NOT the transformation of the bounding box!



M




$$(x'_{max}, y'_{max}, z'_{max}) = (\max(x_0, x_1, x_2, x_3, x_4, x_5, x_6, x_7), \max(y_0, y_1, y_2, y_3, y_4, x_5, x_6, x_7), \max(z_0, z_1, z_2, z_3, z_4, x_5, x_6, x_7))$$

Questions?

Are Bounding Volumes Enough?


- If ray hits bounding volume,
must we test all primitives inside it?
 - Lots of work, think of a 1M-triangle mesh

**bounding
sphere**




Bounding Volume Hierarchies

- If ray hits bounding volume, must we test all primitives inside it?
 - Lots of work, think of a 1M-triangle mesh
- You guessed it already, we'll split the primitives in groups and build recursive bounding volumes
 - Like collision detection, remember?




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Bounding Volume Hierarchy (BVH)

- Find bounding box of objects/primitives
- Split objects/primitives into two, compute child BVs
- Recurse, build a binary tree




Where to Split Objects?

- At midpoint of current volume *OR*
- Sort, and put half of the objects on each side *OR*
- Use modeling hierarchy




Where to Split Objects?

- At midpoint of current volume *OR*
- Sort, and put half of the objects on each side *OR*
- Use modeling hierarchy




Questions?!


Ray-BVH Intersection




Ray-BVH Intersection




Ray-BVH Intersection




Intersection with BVH



Intersection with BVH



Intersection with BVH



BVH Discussion

- Advantages
 - easy to construct
 - easy to traverse
 - binary tree (=simple structure)
- Disadvantages
 - may be difficult to choose a good split for a node
 - poor split may result in minimal spatial pruning


BVH Discussion

- Advantages
 - easy to construct
 - easy to traverse
 - binary tree (=simple structure)
- Disadvantages
 - may be difficult to choose a good split for a node
 - poor split may result in minimal spatial pruning
- Still one of the best methods
 - **Recommended for your first hierarchy!**

- Advantages
 - easy to construct
 - easy to traverse
 - binary tree (=simple structure)
- Disadvantages
 - may be difficult to choose a good split for a node
 - poor split may result in minimal spatial pruning
- Still one of the best methods
 - **Recommended for your first hierarchy!**

Kd-trees


- Probably most popular acceleration structure
- Binary tree, axis-aligned splits
 - Each node splits space in half along an axis-aligned plane
- A **space partition**: The nodes do not overlap!
 - This is in contrast to BVHs



Data Structure

KdTreeNode:

```
KdTreeNode* backNode, frontNode //children  
int dimSplit // either x, y or z  
float splitDistance  
    // from origin along split axis  
boolean isLeaf  
List of triangles //only for leaves
```




here dimSplit = 0 (x axis)

X=splitDistance


Kd-tree Construction

- Start with scene axis-aligned bounding box
- Decide which dimension to split (e.g. longest)
- Decide at which distance to split (not so easy)




Kd-tree Construction - Split

- Distribute primitives to each side
- If a primitive overlaps split plane, assign to both sides



Kd-tree Construction - Recurse

- Stop when minimum number of primitives reached
- Other stopping criteria possible



Questions?


- Further reading on efficient Kd-tree construction
 - [Hunt, Mark & Stoll, IRT 2006](#)
 - [Zhou et al., SIGGRAPH Asia 2008](#)

Zhou et al.




Kd-tree Traversal - High Level

- If leaf, intersect with list of primitives
- If intersects back child, recurse
- If intersects front child, recurse




Kd-tree Traversal, Naïve Version

- Could use bounding box test for each child
- But redundant calculation: bbox similar to that of parent node, plus axis aligned, one single split




Kd-tree Traversal, Smarter Version

- Get main bbox intersection from parent
 - t_{near} , t_{far}
- Intersect with splitting plane
 - easy because axis aligned



Kd-tree Traversal - Three Cases


- Intersects only back, only front, or both
- Can be tested by examining t , t_{start} and t_{end}



Kd-tree traversal - three cases

- If $t > t_{\text{end}} \Rightarrow$ intersect only front
- If $t < t_{\text{start}} \Rightarrow$ intersect only back

Note: “Back” and “Front” depend on ray direction!



Kd-tree Traversal Pseudocode

```
travers(orig, dir, t_start, t_end):
```

```
    #adapted from Ingo Wald's thesis
```

```
    #assumes that dir[self.dimSplit] > 0
```

```
    if self.isLeaf:
```

```
        return intersect(self.listOfTriangles, orig, dir, t_start, t_end)
```

```
    t = (self.splitDist - orig[self.dimSplit]) / dir[self.dimSplit];
```

```
    if t <= t_start:
```

```
        # case one, t <= t_start <= t_end -> cull front side
```

```
        return self.backSideNode.traverse(orig, dir, t_start, t_end)
```

```
    elif t >= t_end:
```

```
        # case two, t_start <= t_end <= t -> cull back side
```

```
        return self.frontSideNode.traverse(orig, dir, t_start, t_end)
```




```
    else:
```

```
        # case three: traverse both sides in turn
```

```
        t_hit = self.frontSideNode.traverse(orig, dir, t_start, t)
```

```
        if t_hit <= t: return t_hit; # early ray termination
```

```
        return self.backSideNode.traverse(orig, dir, t, t_end)
```



Important!

```
travers(orig, dir, t_start, t_end):
```

```
    #adapted from Ingo Wald's thesis
```

```
    #assumes that dir[self.dimSplit] > 0
```

```
    if self.isLeaf:
```

```
        return intersect(self.listOfTriangles, orig, dir, t_start, t_end)
```

```
    t = (self.splitDist - orig[self.dimSplit]) / dir[self.dimSplit];
```

```
    if t <= t_start:
```

```
        # case one, t <= t_start <= t_end -> cull front side
```

```
        return self.backSideNode.traverse(orig, dir, t_start, t_end)
```

```
    elif t >= t_end:
```

```
        # case two, t_start <= t_end <= t -> cull back side
```

```
        return self.frontSideNode.traverse(orig, dir, t_start, t_end)
```




```
    else:
```

```
        # case three: traverse both sides in turn
```

```
        t_hit = self.frontSideNode.traverse(orig, dir, t_start, t)
```




```
        if t_hit <= t: return t_hit; # early ray termination
```

```
        return self.backSideNode.traverse(orig, dir, t, t_end)
```




Early termination is powerful!

```
travers(orig, dir, t_start, t_end):  
    #adapted from Ingo Wald's thesis  
    #assumes that dir[self.dimSplit] > 0  
    if self.isLeaf:  
        return intersect(self.listOfTriangles, orig, dir, t_start, t_end)  
    t = (self.splitDist - orig[self.dimSplit]) / dir[self.dimSplit];  
    if t <= t_start:  
        # case one, t <= t_start <= t_end -> cull front side  
        return self.backSideNode.traverse(orig, dir, t_start, t_end)  
    elif t >= t_end:  
        # case two, t_start <= t_end <= t -> cull back side  
        return self.frontSideNode.traverse(orig, dir, t_start, t_end)  
    else:  
        # case three: traverse both sides in turn  
        t_hit = self.frontSideNode.traverse(orig, dir, t_start, t)  
        if t_hit <= t: return t_hit; # early ray termination  
        return self.backSideNode.traverse(orig, dir, t, t_end)
```



Early termination is powerful

- If there is an intersection in the first node, don't visit the second one
- Allows ray casting to be reasonably independent of scene depth complexity



Recap: Two main gains

- Only intersect with triangles “near” the line
- Stop at the first intersection

Two main gains

```
travers(orig, dir, t_start, t_end):
```

```
    #adapted from Ingo Wald's thesis
```

```
    #assumes that dir[self.dimSplit] > 0
```

```
    if self.isLeaf:
```

```
        return intersect(self.listOfTriangles, orig, dir, t_start, t_end)
```

```
    t = (self.splitDist - orig[self.dimSplit]) / dir[self.dimSplit];
```

```
    if t <= t_start:
```

```
        # case one, t <= t_start <= t_end -> cull front side
```

```
        return self.backSideNode.traverse(orig, dir, t_start, t_end)
```

```
    elif t >= t_end:
```

```
        # case two, t_start <= t_end <= t -> cull back side
```

```
        return self.frontSideNode.traverse(orig, dir, t_start, t_end)
```

```
    else:
```




```
        # case three: traverse both sides in turn
```

```
        t_hit = self.frontSideNode.traverse(orig, dir, t_start, t)
```

```
        if t_hit <= t: return t_hit; # early ray termination
```

```
        return self.backSideNode.traverse(orig, dir, t, t_end)
```

Only near line



stop at first intersection


Important Details

- For leaves, do NOT report intersection if t is not in $[t_{\text{near}}, t_{\text{far}}]$.
 - Important for primitives that overlap multiple nodes!
- Need to take direction of ray into account
 - Reverse back and front if the direction has negative coordinate along the split dimension
- Degeneracies when ray direction is parallel to one axis

- For leaves, do NOT report intersection if t is not in $[t_{\text{near}}, t_{\text{far}}]$.
 - Important for primitives that overlap multiple nodes!
- Need to take direction of ray into account
 - Reverse back and front if the direction has negative coordinate along the split dimension
- Degeneracies when ray direction is parallel to one axis


Where to split for construction?

- Example for baseline
- Note how this ray traverses easily: one leaf only




Split in the Middle

- Does not conform to empty vs. dense areas
- Inefficient traversal – Not so good!



Split in the Median

- Tries to balance tree, but does not conform to empty vs. dense areas
- Inefficient traversal – Not good




Optimizing Splitting Planes

- Most people use the Surface Area Heuristic (SAH)
 - [MacDonald and Booth 1990, “Heuristic for ray tracing using space subdivision”, Visual Computer](#)
- Idea: simple probabilistic prediction of traversal cost based on split distance
- Then try different possible splits and keep the one with lowest cost
- Further reading on efficient Kd-tree construction
 - [Hunt, Mark & Stoll, IRT 2006](#)
 - [Zhou et al., SIGGRAPH Asia 2008](#)

Surface Area Heuristic

- Probability that we need to intersect a child
 - Area of the bbox of that child
(exact for uniformly distributed rays)
- Cost of the traversal of that child
 - number of primitives (simplistic heuristic)
- This heuristic likes to put big densities of primitives in small-area nodes



Is it Important to Optimize Splits?

- Given the same traversal code, the quality of Kd-tree construction can have a big impact on performance, e.g. a factor of 2 compared to naive middle split
 - But then, you should consider carefully if you need that extra performance
 - Could you optimize something else for bigger gain?

Hard-core efficiency considerations

- See e.g. Ingo Wald's PhD thesis
 - <http://www.mpi-inf.mpg.de/~wald/PhD/>
- Calculation
 - Optimized barycentric ray-triangle intersection
- Memory
 - Make kd-tree node as small as possible (dirty bit packing, make it 8 bytes)
- Parallelism
 - SIMD extensions, trace 4 rays at a time, mask results where they disagree

Pros and Cons of Kd trees

- Pros
 - Simple code
 - Efficient traversal
 - Can conform to data
- Cons
 - costly construction, not great if you work with moving objects

Questions?

- For extensions to moving scenes, see [Real-Time KD-Tree Construction on Graphics Hardware, Zhou et al., SIGGRAPH 2008](#)



Stack Studios, Rendered using [Maxwell](#)

Questions?