


Graphics Pipeline & Rasterization II

calvin
and
Hobbes





WATERSON



MIT EECS 6.837
Computer Graphics
Wojciech Matusik



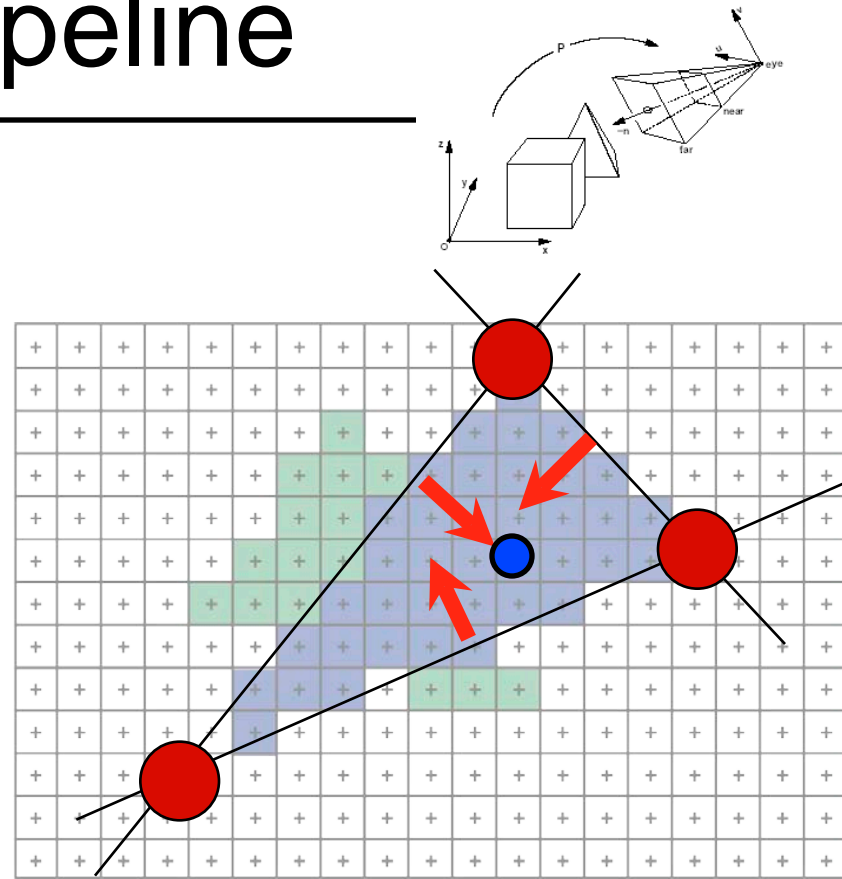
Modern Graphics Pipeline

- Project vertices to 2D (image)
- Rasterize triangle: find which pixels should be lit
- Compute per-pixel color
- Test visibility (Z-buffer), update frame buffer color







Modern Graphics Pipeline

- Project vertices to 2D (image)
- Rasterize triangle: find which pixels should be lit
 - For each pixel, test 3 edge equations
 - if all pass, draw pixel
- Compute per-pixel color
- Test visibility (Z-buffer), update frame buffer color




Modern Graphics Pipeline

- Perform projection of vertices
- Rasterize triangle: find which pixels should be lit
- Compute per-pixel color
- Test visibility,
update frame buffer color
 - Store minimum distance to camera for each pixel in “Z-buffer”
 - ~same as t_{\min} in ray casting!
 - if $\text{new_z} < \text{zbuffer}[x,y]$
 $\text{zbuffer}[x,y] = \text{new_z}$
 $\text{framebuffer}[x,y] = \text{new_color}$








frame buffer



Z buffer






Modern Graphics Pipeline

```
For each triangle
  transform into eye space
  (perform projection)
  setup 3 edge equations
  for each pixel x,y
    if passes all edge equations
      compute z
      if  $z < \text{zbuffer}[x,y]$ 
         $\text{zbuffer}[x,y] = z$ 
         $\text{framebuffer}[x,y] = \text{shade}()$ 
```



Modern Graphics Pipeline


```
For each triangle
  transform into eye space
  (perform projection)
  setup 3 edge equations
  for each pixel x,y
    if passes all edge equations
      compute z
      if  $z < \text{zbuffer}[x,y]$ 
         $\text{zbuffer}[x,y] = z$ 
         $\text{framebuffer}[x,y] = \text{shade}()$ 
```



Questions?


Visibility

- How do we know which parts are visible/in front?




Ray Casting

- Maintain intersection with closest object




Visibility

- In ray casting, use intersection with closest t
- Now we have swapped the loops (pixel, object)
- What do we do?





Z buffer

- In addition to frame buffer (R, G, B)
- Store distance to camera (z -buffer)
- Pixel is updated only if $newz$ is closer than z -buffer value




Works for hard cases!



Interpolation in Screen Space


- How do we get that Z value for each pixel?
 - We only know z at the vertices...
 - (Remember, screen-space z is actually z'/w')
 - Must interpolate from vertices into triangle interior

```
For each triangle
  for each pixel (x,y)
    if passes all edge equations
      compute z
      if z < zbuffer[x,y]
        zbuffer[x,y] = z
        framebuffer[x,y] = shade()
```




Interpolation in Screen Space


- Also need to interpolate color, normals, texture coordinates, etc. between vertices
 - We did this with barycentrics in ray casting
 - Linear interpolation in object space
 - Is this the same as linear interpolation on the screen?



Interpolation in Screen Space




Interpolation in Screen Space



Nope, Not the Same

- Linear variation in world space does not yield linear variation in screen space due to projection
 - Think of looking at a checkerboard at a steep angle; all squares are the same size on the plane, but not on screen

Image: Wikipedia




Head-on view



linear screen-space
("Gouraud") interpolation

BAD



Perspective-correct
Interpolation


Back to the basics: Barycentrics

- Barycentric coordinates for a triangle (**a**, **b**, **c**)

$$P(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$$


– Remember, $\alpha + \beta + \gamma = 1$, $\alpha, \beta, \gamma \geq 0$

- Barycentrics are very general:
 - Work for x, y, z, u, v, r, g, b
 - Anything that varies linearly in object space
 - including z




Basic strategy

- Given screen-space x', y'
- Compute barycentric coordinates
- Interpolate anything specified at the three vertices



Basic strategy

- How to make it work
 - start by computing x', y' given barycentrics
 - invert
- Later: shortcut barycentrics, directly build interpolants



From barycentric to screen-space

- Barycentric coordinates for a triangle (**a**, **b**, **c**)

$$P(\alpha, \beta, \gamma) = \alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}$$

– Remember, $\alpha + \beta + \gamma = 1$, $\alpha, \beta, \gamma \geq 0$

- Let's project point P by projection matrix **C**

$$\begin{aligned} \mathbf{C}P &= \mathbf{C}(\alpha \mathbf{a} + \beta \mathbf{b} + \gamma \mathbf{c}) \\ &= \alpha \mathbf{C}\mathbf{a} + \beta \mathbf{C}\mathbf{b} + \gamma \mathbf{C}\mathbf{c} \\ &= \alpha \mathbf{a}' + \beta \mathbf{b}' + \gamma \mathbf{c}' \end{aligned}$$

a', **b'**, **c'** are the
projected
homogeneous
vertices before
division by w

Projection

- Let's use simple formulation of projection going from 3D homogeneous coordinates to 2D homogeneous coordinates

$$C = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$

- No crazy near-far or storage of $1/z$
- We use ' for screen space coordinates

From barycentric to screen-space

- From previous slides:

$$P' = CP = \alpha \mathbf{a}' + \beta \mathbf{b}' + \gamma \mathbf{c}'$$

\mathbf{a}' , \mathbf{b}' , \mathbf{c}' are the
projected
homogeneous
vertices

- Seems to suggest it's linear in screen space.
But it's homogenous coordinates

From barycentric to screen-space

- From previous slides:

$$P' = CP = \alpha \mathbf{a}' + \beta \mathbf{b}' + \gamma \mathbf{c}'$$


\mathbf{a}' , \mathbf{b}' , \mathbf{c}' are the
projected
homogeneous
vertices

- Seems to suggest it's linear in screen space.
But it's homogenous coordinates
- After division by w , the (x,y) screen coordinates are

$$(P'_x/P'_w, P'_y/P'_w) = \left(\frac{\alpha a'_x + \beta b'_x + \gamma c'_x}{\alpha a'_w + \beta b'_w + \gamma c'_w}, \frac{\alpha a'_y + \beta b'_y + \gamma c'_y}{\alpha a'_w + \beta b'_w + \gamma c'_w} \right)$$

Recap: barycentric to screen-space

$$\begin{array}{c} \text{projective} \\ \text{equivalence} \\ \text{(up to scale)} \end{array} \begin{array}{c} \text{2D} \\ \text{homogenous} \\ \text{coordinates} \end{array} \begin{array}{c} \text{projected} \\ \text{vertices} \end{array} \begin{array}{c} \text{barycentric} \\ \text{coordinates} \end{array}$$
$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} \sim \begin{pmatrix} P'_x \\ P'_y \\ P'_w \end{pmatrix} = \begin{pmatrix} a'_x & b'_x & c'_x \\ a'_y & b'_y & c'_y \\ a'_z & b'_z & c'_z \end{pmatrix} \begin{pmatrix} \alpha \\ \beta \\ \gamma \end{pmatrix}$$



From screen-space to barycentric

$$\begin{array}{c} \text{projective} \\ \text{equivalence} \\ \text{(up to scale)} \end{array} \begin{array}{c} \text{2D} \\ \text{homogenous} \\ \text{coordinates} \end{array} \begin{array}{c} \text{projected} \\ \text{vertices} \end{array} \begin{array}{c} \text{barycentric} \\ \text{coordinates} \end{array}$$
$$\begin{pmatrix} x' \\ y' \\ 1 \end{pmatrix} \sim \begin{pmatrix} P'_x \\ P'_y \\ P'_w \end{pmatrix} = \begin{pmatrix} a'_x & b'_x & c'_x \\ a'_y & b'_y & c'_y \\ a'_z & b'_z & c'_z \end{pmatrix} \begin{pmatrix} \alpha \\ \beta \\ \gamma \end{pmatrix}$$

- It's a projective mapping from the barycentrics onto screen coordinates!
 - Represented by a 3x3 matrix
- We'll take the inverse mapping to get from (x, y, 1) to the barycentrics!

From Screen to Barycentrics

projective
equivalence

$$\begin{pmatrix} \alpha \\ \beta \\ \gamma \end{pmatrix} \sim \begin{pmatrix} a'_x & b'_x & c'_x \\ a'_y & b'_y & c'_y \\ a'_w & b'_w & c'_w \end{pmatrix}^{-1} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

- Recipe
 - Compute projected homogeneous coordinates \mathbf{a}' , \mathbf{b}' , \mathbf{c}'
 - Put them in the columns of a matrix, invert it
 - Multiply screen coordinates $(x, y, 1)$ by inverse matrix
 - **Then divide by the sum of the resulting coordinates**
 - This ensures the result sums to one like barycentrics should
 - Then interpolate value (e.g. Z) from vertices using them!

From Screen to Barycentrics

$$\begin{pmatrix} \alpha \\ \beta \\ \gamma \end{pmatrix} \sim \begin{pmatrix} a'_x & b'_x & c'_x \\ a'_y & b'_y & c'_y \\ a'_w & b'_w & c'_w \end{pmatrix}^{-1} \begin{pmatrix} x \\ y \\ 1 \end{pmatrix}$$

- Notes:
 - matrix is inverted once per triangle
 - can be used to interpolate z, color, texture coordinates, etc.

Pseudocode – Rasterization

For every triangle

 ComputeProjection

Compute interpolation matrix

 Compute bbox, clip bbox to screen limits

 For all pixels x, y in bbox

 Test edge functions

 If all $E_i > 0$

compute barycentrics

 interpolate z from vertices

 if $z < zbuffer[x, y]$

 interpolate UV coordinates from vertices

 look up texture color k_d

 Framebuffer[x, y] = k_d //or more complex shader



Pseudocode – Rasterization

For every triangle

 ComputeProjection

Compute interpolation matrix

 Compute bbox, clip bbox to screen limits

 For all pixels x, y in bbox

 Test edge functions

 If all $E_i > 0$

compute barycentrics

 interpolate z from vertices

 if $z < zbuffer[x, y]$

 interpolate UV coordinates from vertices

 look up texture color k_d


 Framebuffer[x, y] = k_d //or more complex shader




Questions?

Supersampling


- Trivial to do with rasterization as well
- Often rates of 2x to 8x
- Requires to compute per-pixel average at the end
- Most effective against edge jaggies
- Usually with jittered sampling
 - pre-computed pattern for a big block of pixels




1 Sample / Pixel



4 Samples / Pixel




16 Samples / Pixel



100 Samples / Pixel

Even this
sampling rate
cannot get rid
of all aliasing
artifacts!

**We are really
only pushing
the problem
farther.**




Related Idea: Multisampling

- Problem
 - Shading is very expensive today (complicated shaders)
 - Full supersampling has linear cost in #samples ($k*k$)
- Goal: High-quality edge antialiasing at lower cost
- Solution
 - Compute shading only once per pixel for each primitive, but resolve visibility at “sub-pixel” level
 - Store ($k*\text{width}$, $k*\text{height}$) frame and z buffers, but share shading results between sub-pixels within a real pixel
 - When visibility samples within a pixel hit different primitives, we get an average of their colors
 - Edges get antialiased without large shading cost

Multisampling, Visually


○ = sub-pixel visibility sample



One pixel


Multisampling, Visually

○ = sub-pixel visibility sample



Multisampling, Visually

○ = sub-pixel visibility sample




One pixel

The color is only computed **once per pixel per triangle** and reused for all the visibility samples that are covered by the triangle.

Supersampling, Visually

○ = sub-pixel visibility sample



One pixel

When supersampling, we compute colors independently for all the visibility samples.

Multisampling Pseudocode

```
For each triangle
  For each pixel
    if pixel overlaps triangle
      color=shade() // only once per pixel!
      for each sub-pixel sample
        compute edge equations & z
        if subsample passes edge equations
          && z < zbuffer[subsample]
            zbuffer[subsample]=z
            framebuffer[subsample]=color
```


Multisampling Pseudocode

```
For each triangle
  For each pixel
    if pixel overlaps triangle
      color=shade() // only once per pixel!
      for each sub-pixel sample
        compute edge equations & z
        if subsample passes edge equations
          && z < zbuffer[subsample]
            zbuffer[subsample]=z
            framebuffer[subsample]=color
At display time: //this is called "resolving"
  For each pixel
    color = average of subsamples
```

Multisampling vs. Supersampling

- Supersampling
 - Compute an entire image at a higher resolution, then downsample (blur + resample at lower res)
- Multisampling
 - Supersample visibility, compute expensive shading only once per pixel, reuse shading across visibility samples
- But Why?
 - Visibility edges are where supersampling really works
 - Shading can be prefiltered more easily than visibility
- This is how GPUs perform antialiasing these days


Multisampling vs. Supersampling

- Supersampling
 - Compute an entire image at a higher resolution, then downsample (blur + resample at lower res)
- Multisampling
 - Supersample visibility, compute expensive shading only once per pixel, reuse shading across visibility samples
- But Why?
 - Visibility edges are where supersampling really works
 - Shading can be prefiltered more easily than visibility
- This is how GPUs perform antialiasing these days


Questions?

Examples of Texture Aliasing

Magnification




Minification




Texture Filtering

- Problem: Prefiltering is impossible when you can only take point samples
 - This is why visibility (edges) need supersampling
- Texture mapping is simpler
 - Imagine again we are looking at an infinite textured plane



Texture Filtering


- We should pre-filter image function *before sampling*
 - That means blurring the image function with a low-pass filter (convolution of image function and filter)



Texture Filtering


- We can combine low-pass and sampling
 - The value of a sample is the integral of the product of the image f and the filter h centered at the sample location
 - “A local average of the image f weighted by the filter h ”

$$\hat{f}_i = \int_{\text{image}} f(x) h(x) dx$$



Texture Filtering


- Well, we can just as well change variables and compute this integral *on the textured plane instead*
 - In effect, we are projecting the pre-filter onto the plane



Texture Filtering


- Well, we can just as well change variables and compute this integral *on the textured plane instead*
 - In effect, we are projecting the pre-filter onto the plane
 - It's still a weighted average of the texture under filter

$$\hat{f}_i = \int_{\text{plane}} f(x') h(x') |J(x, x')| dx'$$



Low-pass filter


Texture Pre-Filtering, Visually



- Must still integrate product of projected filter and texture


Solution: Precomputation

- We'll precompute and store a set of prefiltered results from each texture with different sizes of prefilters




Solution: Precomputation

- We'll precompute and store a set of prefiltered results from each texture with different sizes of prefilters




Solution: Precomputation

- We'll precompute and store a set of prefiltered results from each texture with different sizes of prefilters
 - Because it's low-passed, we can also subsample




Solution: Precomputation

- We'll precompute and store a set of prefiltered results from each texture with different sizes of prefilters
 - Because it's low-passed, we can also subsample





This is Called “MIP-Mapping”

- Construct a pyramid of images that are pre-filtered and re-sampled at $1/2$, $1/4$, $1/8$, etc., of the original image's sampling
- During rasterization we compute the index of the decimated image that is sampled at a rate closest to the density of our desired sampling rate
- MIP stands for *multum in parvo* which means *many in a small place*



Storing MIP Maps


- Can be stored compactly: Only 1/3 more space!



MIP-Mapping

- When a pixel wants an integral of the pre-filtered texture, we must find the “closest” results from the precomputed MIP-map pyramid
 - Must compute the “size” of the projected pre-filter in the texture UV domain

Projected pre-filter




MIP-Mapping

- Simplest method: Pick the scale closest, then do usual reconstruction on that level (e.g. bilinear between 4 closest texture pixels)




Projected pre-filter



MIP-Mapping

- Simplest method: Pick the scale closest, then do usual reconstruction on that level (e.g. bilinear between 4 closest texture pixels)
- Problem: discontinuity when switching scale



Projected pre-filter



Tri-Linear MIP-Mapping


- Use **two** closest scales, compute reconstruction results from both, and linearly interpolate between them

Projected pre-filter




Tri-Linear MIP-Mapping

- Use **two** closest scales, compute reconstruction results from both, and linearly interpolate between them
- Problem: our filter might not be circular, because of foreshortening




Projected pre-filter




Anisotropic filtering

- Approximate Elliptical filter with multiple circular ones (usually 5)
- Perform trilinear lookup at each one
- i.e. consider five times eight values
 - fair amount of computation
 - this is why graphics hardware has dedicated units to compute trilinear mipmap reconstruction


Projected pre-filter



MIP Mapping Example




Nearest Neighbor




MIP Mapped (Tri-Linear)

MIP Mapping Example



MIP Mapping Example


Questions



Finding the MIP Level

- Often we think of the pre-filter as a box
 - What is the projection of the square pixel “window” in texture space?


Projected pre-filter



Finding the MIP Level

- Often we think of the pre-filter as a box
 - What is the projection of the square pixel “window” in texture space?
 - Answer is in the partial derivatives p_x and p_y of (u,v) w.r.t. screen (x,y)


- **Projection of pixel center**
Projected pre-filter



For isotropic trilinear mipmapping

- No right answer, circular approximation
- Two most common approaches are
 - Pick level according to the length (in texels) of the longer partial
 $\log_2 \max \{w|p_x|, h|p_y|\}$
 - Pick level according to the length of their sum
 $\log_2 \sqrt{(w|p_x|)^2 + (h|p_y|)^2}$


- **Projection of pixel center**
Projected pre-filter



Anisotropic filtering

- Pick levels according to smallest partial
 - well, actually max of the smallest and the largest/5
- Distribute circular “probes” along longest one
- Weight them by a Gaussian

- **Projection of pixel center**
Projected pre-filter



How Are Partials Computed?

- You can derive closed form formulas based on the uv and xyw coordinates of the vertices...
 - This is what used to be done
- ..but shaders may compute texture coordinates programmatically, not necessarily interpolated
 - No way of getting analytic derivatives!
- **In practice, use finite differences**
 - GPUs process pixels in blocks of (at least) 4 anyway
 - These 2x2 blocks are called *quads*






Image Quality Comparison



trilinear mipmapping
(excessive blurring)



anisotropic filtering

Further Reading

- Paul Heckbert published seminal work on texture mapping and filtering in his master's thesis (!)
 - Including EWA
 - Highly recommended reading!
 - See <http://www.cs.cmu.edu/~ph/textfund/textfund.pdf>
- More reading
 - [Feline: Fast Elliptical Lines for Anisotropic Texture Mapping, McCormack, Perry, Farkas, Jouppe SIGGRAPH 1999](#)
 - [Texram: A Smart Memory for Texturing Schilling, Knittel, Strasser, IEEE CG&A, 16\(3\): 32-41](#)

Arf!



- Paul Heckbert published seminal work on texture mapping and filtering in his master's thesis (!)
 - Including EWA
 - Highly recommended reading!
 - See <http://www.cs.cmu.edu/~ph/textfund/textfund.pdf>
- More reading
 - [Feline: Fast Elliptical Lines for Anisotropic Texture Mapping, McCormack, Perry, Farkas, Jouppe SIGGRAPH 1999](#)
 - [Texram: A Smart Memory for Texturing Schilling, Knittel, Strasser, IEEE CG&A, 16\(3\): 32-41](#)

Arf!



Ray Casting vs. Rasterization

Ray Casting

For each pixel

For each object

- Ray-centric
- Needs to store scene in memory
- (Mostly) Random access to scene

Rasterization

For each triangle

For each pixel

- Triangle centric
- Needs to store image (and depth) into memory
- (Mostly) random access to frame buffer

Which is smaller? Scene or Frame?

Frame

Which is easiest to access randomly?

Frame because regular sampling


Good References

- <http://www.tomshardware.com/reviews/ray-tracing-rasterization,2351.html>
- <http://c0de517e.blogspot.com/2011/09/raytracing-myths.html>
- <http://people.csail.mit.edu/fredo/tmp/rendering.pdf>

Graphics Hardware


- High performance through
 - Parallelism
 - Specialization
 - No data dependency
 - Efficient pre-fetching
- More next week

task
parallelism




- High performance through
 - Parallelism
 - Specialization
 - No data dependency
 - Efficient pre-fetching
- More next week

task
parallelism



Movies

both rasterization and ray tracing



Games

rasterization




CAD-CAM & Design

*rasterization for GUI,
anything for final image*




Architecture

*ray-tracing, rasterization with
preprocessing for complex lighting*






Virtual Reality

rasterization







Visualization

*mostly rasterization,
interactive ray-tracing is starting*



Medical Imaging

*same as
visualization*



Questions?
